

Agentic AI **guardrails** are the rules and safety checks that guide AI systems to act responsibly, ensuring their decisions stay aligned with human goals, values, and safety limits.

